

# MADMO

## Boostings and ensembles. Part 2.

Taras Khakhulin

Deep Learning Engineer Samsung AI Center

Skoltech & MIPT MS

[t.khakhulin@gmail.com](mailto:t.khakhulin@gmail.com)

<https://github.com/khakhulin/>

[https://twitter.com/t\\_khakhulin](https://twitter.com/t_khakhulin)

# Ensemble problems

Fixed target function

Similar dataset

Solve one task

# Ensembles

## Voting

- averaging

## Stacking

## Boosting

## Output Coding

- code target (squared)


## Bagging

## Heuristics

- Hand-crafted methods

# Random Forest


Bagging + RSM = Random Forest



# Lecture 1. Part 2

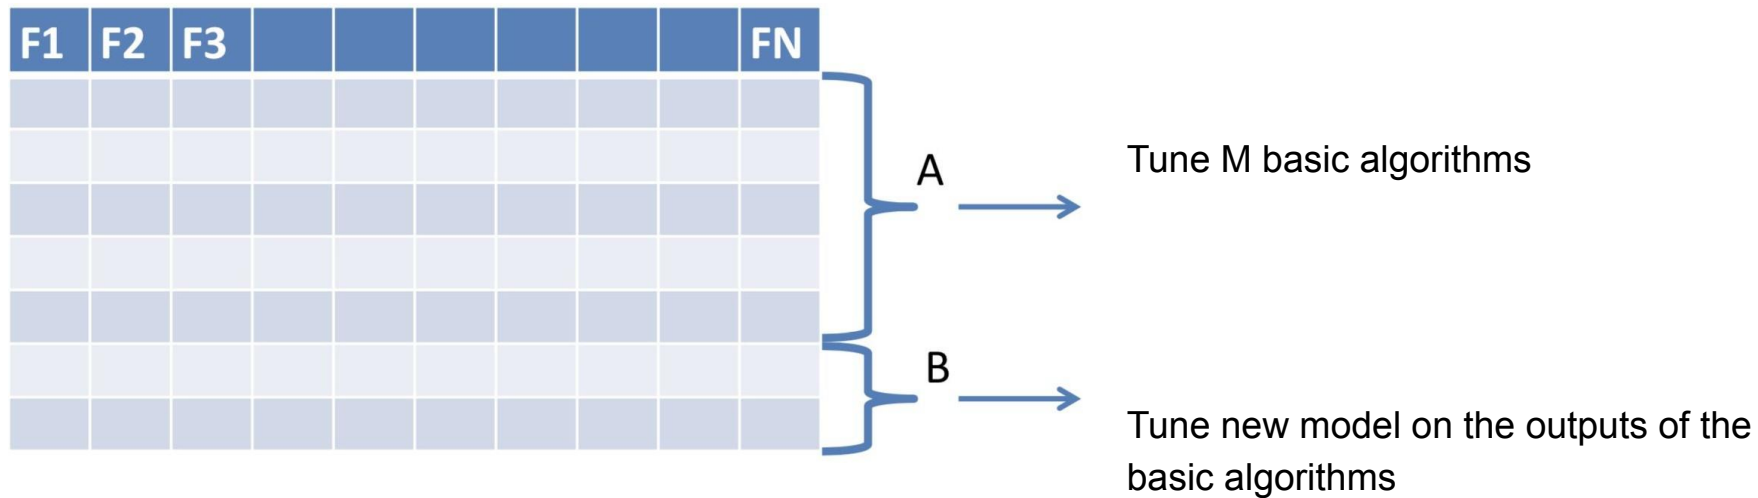
# Stacking

How to build an ensemble from *different* models?



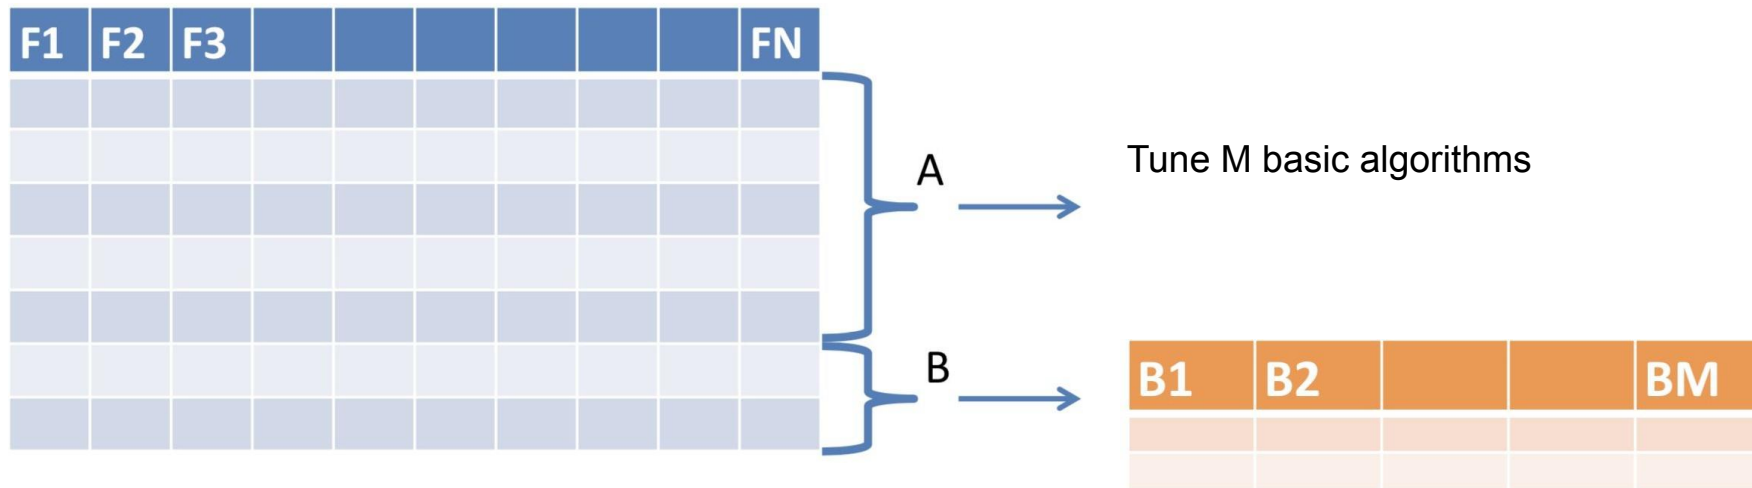
# Stacking

How to build an ensemble from *different* models?



# Stacking


How to build an ensemble from *different* models?





# Stacking

How to build an ensemble from *different* models?



$$a(x) = \sum_{t=1}^T \alpha_t b_t(x)$$

e.g.

# Stacking

How to build an ensemble from *different* models?

# Stacking

How to build an ensemble from *different* models?

- Use different datasets (or datasets parts) for different level models.

# Stacking

How to build an ensemble from *different* models?

- Use different datasets (or datasets parts) for different level models.
- Experiment with different models (linear, trees ensembles, simple networks, etc.)

# Stacking

How to build an ensemble from *different* models?

- Use different datasets (or datasets parts) for different level models.
- Experiment with different models (linear, trees ensembles, simple networks, etc.)
- Or just different GBT ensembles (hola, kaggle :)

# Blending

Just combine several *strong/complex* models.

Weights should sum up to 1  
and come from  $[0; 1]$

$$a(x) = \sum_{t=1}^T \alpha_t b_t(x)$$

# Blending

Just combine several *strong/complex* models.

Weights should sum up to 1  
and come from  $[0; 1]$

$$a(x) = \sum_{t=1}^T \alpha_t b_t(x)$$

- Simple and intuitive ensembling method

# Blending

Just combine several *strong/complex* models.

Weights should sum up to 1  
and come from  $[0; 1]$

$$a(x) = \sum_{t=1}^T \alpha_t b_t(x)$$

- Simple and intuitive ensembling method
- Finding optimal weights could be tricky



# Blending

Just combine several *strong/complex* models.

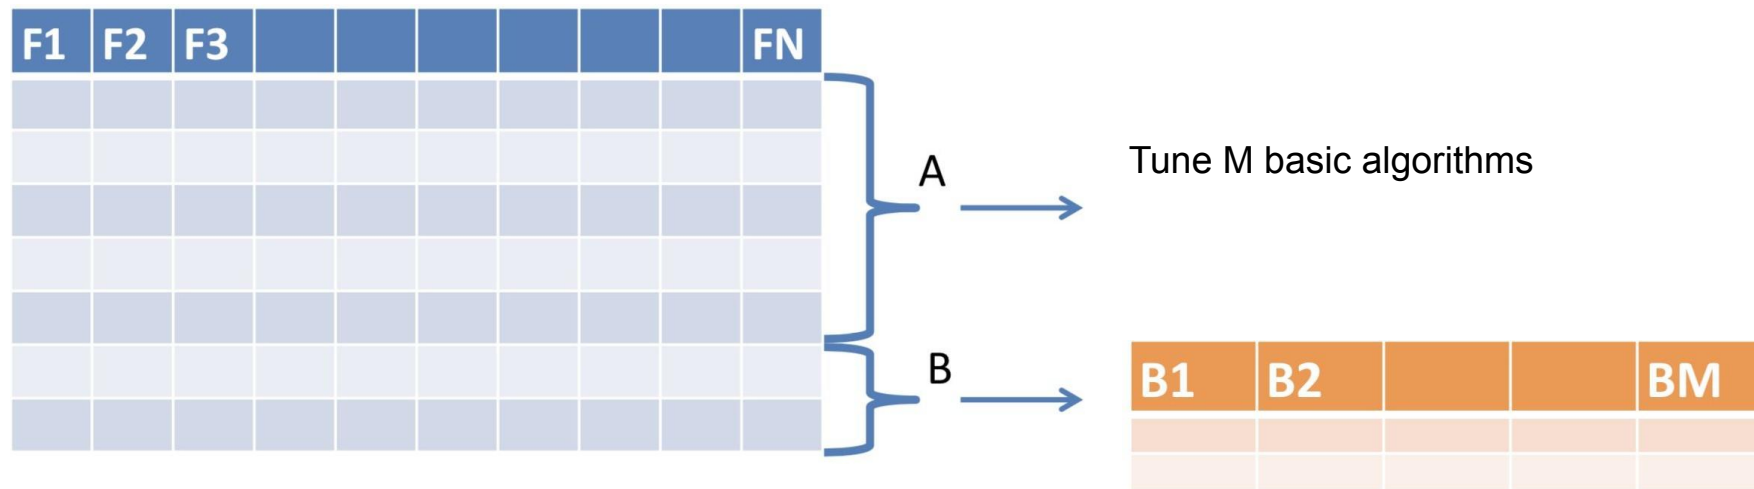
Weights should sum up to 1  
and come from  $[0; 1]$

$$a(x) = \sum_{t=1}^T \alpha_t b_t(x)$$

- Simple and intuitive ensembling method
- Finding optimal weights could be tricky
- Linear composition is not always enough

# Blending as a Stacking


How to build an ensemble from *different* models?




$$a(x) = \sum_{t=1}^T \alpha_t b_t(x)$$

e.g.

# Blending as a Stacking

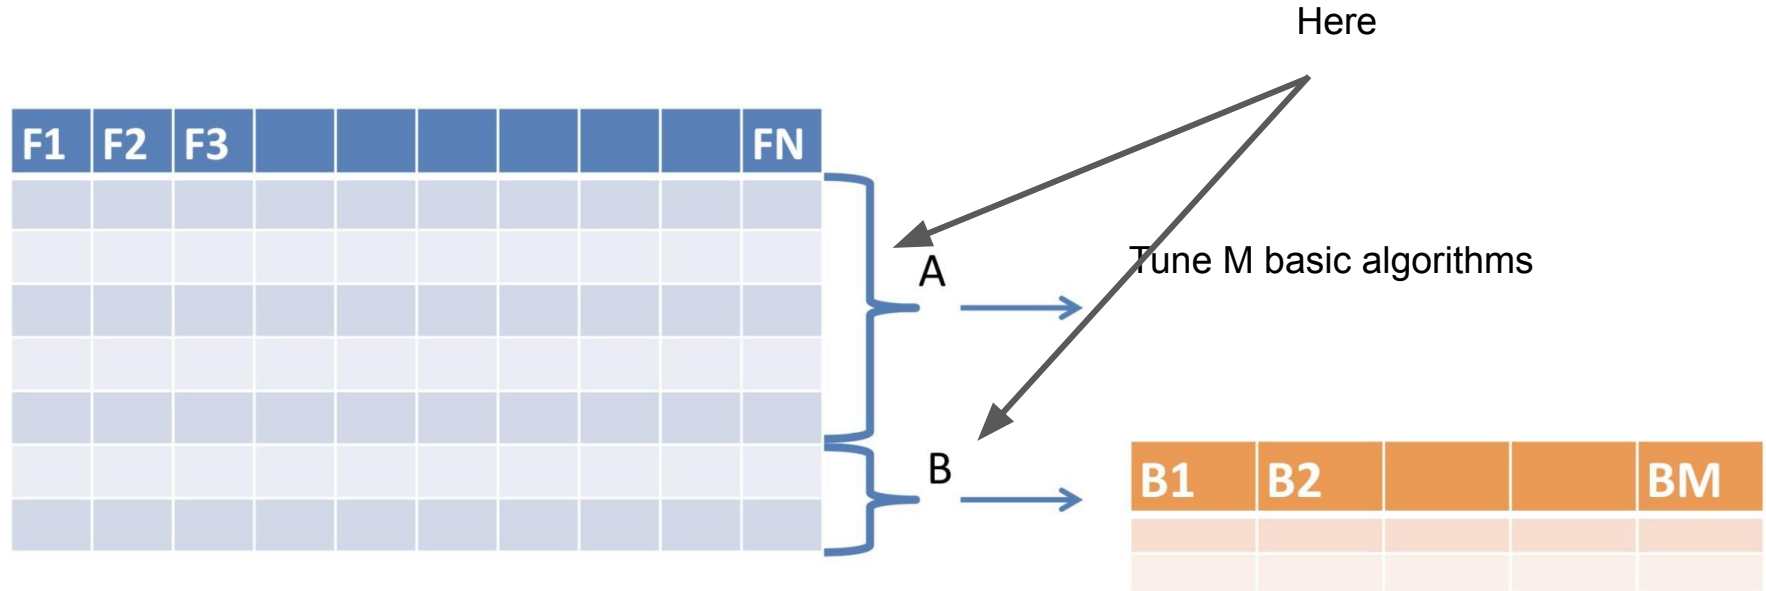


# Blending as a Stacking




Where is a problem?


# Blending as a Stacking



# Blending-2




# Blending-2



2nd meta  
model

1st meta  
model

# Blending-2



2nd meta  
model

Ensemble them

1st meta  
model



# Still confused about Stacking?

A				
X0	x1	x2	xn	y
0.17	0.25	0.93	0.79	1
0.35	0.61	0.93	0.57	0
0.44	0.59	0.56	0.46	0
0.37	0.43	0.74	0.28	1
0.96	0.07	0.57	0.01	1

B				
X0	x1	x2	xn	y
0.89	0.72	0.50	0.66	0
0.58	0.71	0.92	0.27	1
0.10	0.35	0.27	0.37	0
0.47	0.68	0.30	0.98	0
0.39	0.53	0.59	0.18	1

C				
X0	x1	x2	xn	y
0.29	0.77	0.05	0.09	?
0.38	0.66	0.42	0.91	?
0.72	0.66	0.92	0.11	?
0.70	0.37	0.91	0.17	?
0.59	0.98	0.93	0.65	?

Train algorithm **0** on A and make predictions for B and C and save to **B1, C1**

Train algorithm **1** on A and make predictions for B and C and save to **B1, C1**

Train algorithm **2** on A and make predictions for B and C and save to **B1, C1**

B1			
pred0	pred1	pred2	y
0.24	0.72	0.70	0
0.95	0.25	0.22	1
0.64	0.80	0.96	0
0.89	0.58	0.52	0
0.11	0.20	0.93	1

C1				
pred0	pred1	pred2	y	Preds3
0.50	0.50	0.39	?	0.45
0.62	0.59	0.46	?	0.23
0.22	0.31	0.54	?	0.99
0.90	0.47	0.09	?	0.34
0.20	0.09	0.61	?	0.05

Train algorithm **3** on B1 and make predictions for C1

# Stacking

- Correlated models
- Use models with different natures
- Understand your models
- Work with your feature space
- Stack over stacking [StackNet](#)

# Stacking


- Correlated models
- Use models with different natures
- Understand your models
- Work with your feature space
- Stack over stacking [StackNet](#) or <http://ml-ensemble.com/>

```
from mlens.ensemble import SuperLearner
ensemble = SuperLearner()
ensemble.add(estimators)
ensemble.add_meta(meta_estimator)
ensemble.fit(X, y).predict(X)
```

# Gradient boosting

# Gradient boosting


$$a_n(x) = h_1(x) + \cdots + h_n(x)$$





# Boosting: intuition

Binary classification problem.


Models - decision stumps.





# Boosting: intuition





$t = 1$




# Boosting: intuition



$t = 1$





$t = 2$









# Boosting: intuition




$t = 1$



$t = 2$





$t = 3$



# Boosting: intuition

Binary classification problem.  
Models - decision stumps.



# Boosting: core idea

We train each model in the ensemble so that it corrects the error of the previous one

# Boosting: Step by step

# Boosting: Step by step regression

$$\frac{1}{2} \sum_{i=1}^{\ell} (a(x_i) - y_i)^2 \rightarrow \min_a$$

Result  
algorithm:

$$a_N(x) = \sum_{n=1}^N b_n(x)$$

$$b_1(x) := \arg \min_{b \in \mathcal{A}} \frac{1}{2} \sum_{i=1}^{\ell} (b(x_i) - y_i)^2$$

# Boosting: Step by step regression

$$\frac{1}{2} \sum_{i=1}^{\ell} (a(x_i) - y_i)^2 \rightarrow \min_a$$

Result  
algorithm:

$$a_N(x) = \sum_{n=1}^N b_n(x)$$

$$b_1(x) := \arg \min_{b \in \mathcal{A}} \frac{1}{2} \sum_{i=1}^{\ell} (b(x_i) - y_i)^2$$

$$s_i^{(1)} = y_i - b_1(x_i)$$

$$b_2(x) := \arg \min_{b \in \mathcal{A}} \frac{1}{2} \sum_{i=1}^{\ell} (b(x_i) - s_i^{(1)})^2$$

# Boosting: Step by step regression

$$\frac{1}{2} \sum_{i=1}^{\ell} (a(x_i) - y_i)^2 \rightarrow \min_a$$

Result  
algorithm:

$$a_N(x) = \sum_{n=1}^N b_n(x)$$

$$s_i^{(N)} = y_i - \sum_{n=1}^{N-1} b_n(x_i) = y_i - a_{N-1}(x_i), \quad i = 1, \dots, \ell;$$

$$b_N(x) := \arg \min_{b \in \mathcal{A}} \frac{1}{2} \sum_{i=1}^{\ell} (b(x_i) - s_i^{(N)})^2$$

# Boosting: Step by step regression

$$\frac{1}{2} \sum_{i=1}^{\ell} (a(x_i) - y_i)^2 \rightarrow \min_a$$

Result  
algorithm:

$$a_N(x) = \sum_{n=1}^N b_n(x)$$

$$s_i^{(N)} = y_i - \sum_{n=1}^{N-1} b_n(x_i) = y_i - a_{N-1}(x_i), \quad i = 1, \dots, \ell;$$

$$b_N(x) := \arg \min_{b \in \mathcal{A}} \frac{1}{2} \sum_{i=1}^{\ell} (b(x_i) - s_i^{(N)})^2 \quad s_i^{(N)} = y_i - a_{N-1}(x_i) = - \left. \frac{\partial}{\partial z} \frac{1}{2} (z - y_i)^2 \right|_{z=a_{N-1}(x_i)}$$



# Gradient boosting

Denote dataset  $\{(x_i, y_i)\}$ , loss function  $L(y, f)$

# Gradient boosting

Denote dataset  $\{(x_i, y_i)\}$ .

, loss function  $L(y, f)$

$$a_N(x) = \sum_{n=0}^N \gamma_n b_n(x)$$

# Gradient boosting

Denote dataset  $\{(x_i, y_i)\}$ .

, loss function  $L(y, f)$

$$a_N(x) = \sum_{n=0}^N \gamma_n b_n(x)$$

$$b_0(x) = \frac{1}{\ell} \sum_{i=1}^{\ell} y_i$$

# Gradient boosting

Denote dataset  $\{(x_i, y_i)\}$ , loss function  $L(y, f)$

$$a_N(x) = \sum_{n=0}^N \gamma_n b_n(x)$$

$$b_0(x) = \frac{1}{\ell} \sum_{i=1}^{\ell} y_i$$

Find next algo to  
minimize the error:

$$\sum_{i=1}^{\ell} L(y_i, a_{N-1}(x_i) + \gamma_N b_N(x_i)) \rightarrow \min_{b_N, \gamma_N}$$

# Gradient boosting

Denote dataset  $\{(x_i, y_i)\}$ , loss function  $L(y, f)$

$$a_N(x) = \sum_{n=0}^N \gamma_n b_n(x)$$


$$b_0(x) = \frac{1}{\ell} \sum_{i=1}^{\ell} y_i$$

Find nex algo to  
minimize the error:

$$\sum_{i=1}^{\ell} L(y_i, a_{N-1}(x_i) + \gamma_N b_N(x_i)) \rightarrow \min_{b_N, \gamma_N}$$


What if we could use gradient descent in *space of our models*?

# Gradient boosting: theory



What if we could use gradient descent in *space of our models*?

# Gradient boosting: theory



What if we could use gradient descent in *space of our models*?

# Gradient boosting

Denote dataset  $\{(x_i, y_i)\}$ , loss function  $L(y, f)$

$$a_N(x) = \sum_{n=0}^N \gamma_n b_n(x)$$

$$\sum_{i=1}^{\ell} L(y_i, a_{N-1}(x_i) + s_i) \rightarrow \min_{s_1, \dots, s_{\ell}}$$

$$b_0(x) = \frac{1}{\ell} \sum_{i=1}^{\ell} y_i$$

$$s_i = - \left. \frac{\partial L}{\partial z} \right|_{z=a_{N-1}(x_i)}$$



# Gradient boosting

Denote dataset  $\{(x_i, y_i)\}$ , loss function  $L(y, f)$

$$a_N(x) = \sum_{n=0}^N \gamma_n b_n(x)$$

$$\sum_{i=1}^{\ell} L(y_i, a_{N-1}(x_i) + s_i) \rightarrow \min_{s_1, \dots, s_{\ell}}$$

$$b_0(x) = \frac{1}{\ell} \sum_{i=1}^{\ell} y_i$$

$$s_i = - \left. \frac{\partial L}{\partial z} \right|_{z=a_{N-1}(x_i)}$$

For every object:

$$\left( - \left. \frac{\partial L}{\partial z} \right|_{z=a_{N-1}(x_i)} \right)_{i=1}^{\ell} = - \nabla_z \sum_{i=1}^{\ell} L(y_i, z_i) \Big|_{z_i=a_{N-1}(x_i)}$$

# Gradient boosting

Denote dataset  $\{(x_i, y_i)\}$ , loss function  $L(y, f)$

$$a_N(x) = \sum_{n=0}^N \gamma_n b_n(x) \quad \sum_{i=1}^{\ell} L(y_i, a_{N-1}(x_i) + s_i) \rightarrow \min_{s_1, \dots, s_{\ell}}$$

$$b_0(x) = \frac{1}{\ell} \sum_{i=1}^{\ell} y_i \quad s_i = - \left. \frac{\partial L}{\partial z} \right|_{z=a_{N-1}(x_i)}$$

For every object:

$$\left( - \left. \frac{\partial L}{\partial z} \right|_{z=a_{N-1}(x_i)} \right)_{i=1}^{\ell} = - \nabla_z \sum_{i=1}^{\ell} L(y_i, z_i) \Big|_{z_i=a_{N-1}(x_i)}$$

$$\gamma_N = \arg \min_{\gamma \in \mathbb{R}} \sum_{i=1}^{\ell} L(y_i, a_{N-1}(x_i) + \gamma b_N(x_i))$$

# Gradient boosting: sum up

What we need:

- Data.
- Loss function and its gradient.
- Family of algorithms (with constraints on hyperparameters if necessary).
- Number of iterations  $M$ .
- Initial value (GBM by Friedman): constant.

# Gradient boosting: regularization

Add regularizer

$$a_N(x) = a_{N-1}(x) + \eta \gamma_N b_N(x), \quad \eta \in (0, 1]$$

Number of iteration

Weak learners

# Gradient boosting: regularization

Add regularizer  $a_N(x) = a_{N-1}(x) + \eta \gamma_N b_N(x), \quad \eta \in (0, 1]$

Number of iteration

Weak learners

Stochastic Gradient Boosting - train  $b_i$  on subsample of data

# Gradient boosting: Loss Function

Classification

Logistic Loss

$$L(y, z) = \log(1 + \exp(-yz)).$$

$$b_N = \arg \min_{b \in \mathcal{A}} \sum_{i=1}^{\ell} \left( b(x_i) - \frac{y_i}{1 + \exp(y_i a_{N-1}(x_i))} \right)^2$$

# Gradient boosting: Loss Function

## Classification

### Logistic Loss

$$L(y, z) = \log(1 + \exp(-yz)).$$

$$b_N = \arg \min_{b \in \mathcal{A}} \sum_{i=1}^{\ell} \left( b(x_i) - \frac{y_i}{1 + \exp(y_i a_{N-1}(x_i))} \right)^2$$

## Regression

MSE

MAE

Where are Trees?



# Gradient boosting on Trees

$$b_n(x) = \sum_{j=1}^{J_n} b_{nj}[x \in R_j] \quad \begin{array}{l} j = 1, \dots, J_n \text{ number of leafs} \\ R_j \text{ subspace} \end{array}$$

$$a_N(x) = a_{N-1}(x) + \gamma_N \sum_{j=1}^{J_N} b_{Nj}[x \in R_j]$$

# Gradient boosting on Trees

$$b_n(x) = \sum_{j=1}^{J_n} b_{nj}[x \in R_j] \quad \begin{array}{l} j = 1, \dots, J_n \text{ number of leafs} \\ R_j \text{ subspace} \end{array}$$

$$a_N(x) = a_{N-1}(x) + \gamma_N \sum_{j=1}^{J_N} b_{Nj}[x \in R_j] = a_{N-1}(x) + \sum_{j=1}^{J_N} \gamma_N b_{Nj}[x \in R_j].$$

# Gradient boosting on Trees

$$b_n(x) = \sum_{j=1}^{J_n} b_{nj}[x \in R_j] \quad \begin{array}{l} j = 1, \dots, J_n \text{ number of leafs} \\ R_j \text{ subspace} \end{array}$$

$$a_N(x) = a_{N-1}(x) + \gamma_N \sum_{j=1}^{J_N} b_{Nj}[x \in R_j] = a_{N-1}(x) + \sum_{j=1}^{J_N} \gamma_N b_{Nj}[x \in R_j].$$

$$\sum_{i=1}^{\ell} L \left( y_i, a_{N-1}(x_i) + \sum_{j=1}^{J_N} \gamma_{Nj}[x \in R_j] \right) \rightarrow \min_{\{\gamma_{Nj}\}_{j=1}^{J_N}}$$

# Gradient boosting on Trees

$$b_n(x) = \sum_{j=1}^{J_n} b_{nj}[x \in R_j] \quad \begin{array}{l} j = 1, \dots, J_n \text{ number of leafs} \\ R_j \text{ subspace} \end{array}$$

$$a_N(x) = a_{N-1}(x) + \gamma_N \sum_{j=1}^{J_N} b_{Nj}[x \in R_j] = a_{N-1}(x) + \sum_{j=1}^{J_N} \gamma_N b_{Nj}[x \in R_j].$$

$$\sum_{i=1}^{\ell} L \left( y_i, a_{N-1}(x_i) + \sum_{j=1}^{J_N} \gamma_{Nj}[x \in R_j] \right) \rightarrow \min_{\{\gamma_{Nj}\}_{j=1}^{J_N}}$$

$$\gamma_{Nj} = \arg \min_{\gamma} \sum_{x_i \in R_j} L(y_i, a_{N-1}(x_i) + \gamma),$$

# AdaBoost

$$L(y, z) = e^{-yz}$$

$$L(a, X) = \sum_{i=1}^{\ell} \exp \left( -y_i \sum_{n=1}^N \gamma_n b_n(x_i) \right)$$

$$s_i = - \left. \frac{\partial L(y_i, z)}{\partial z} \right|_{z=a_{N-1}(x_i)} = y_i \underbrace{\exp \left( -y_i \sum_{n=1}^{N-1} \gamma_n b_n(x_i) \right)}_{w_i}.$$

# Extreme Gradient Boosting (XGBoost)

XGBoost

*dmlc*  
***XGBoost***

# XGBoost

$$s = \left( - \frac{\partial L}{\partial z} \Big|_{z=a_{N-1}(x_i)} \right)_{i=1}^{\ell} = -\nabla_z \sum_{i=1}^{\ell} L(y_i, z_i) \Big|_{z_i=a_{N-1}(x_i)}$$

$$b_N(x) = \arg \min_{b \in \mathcal{A}} \sum_{i=1}^{\ell} (b(x_i) - s_i)^2$$



# XGBoost

$$s = \left( - \frac{\partial L}{\partial z} \Big|_{z=a_{N-1}(x_i)} \right)_{i=1}^{\ell} = -\nabla_z \sum_{i=1}^{\ell} L(y_i, z_i) \Big|_{z_i=a_{N-1}(x_i)}$$

$$b_N(x) = \arg \min_{b \in \mathcal{A}} \sum_{i=1}^{\ell} (b(x_i) - s_i)^2$$

Why?

# XGBoost

$$s = \left( - \frac{\partial L}{\partial z} \Big|_{z=a_{N-1}(x_i)} \right)_{i=1}^{\ell} = -\nabla_z \sum_{i=1}^{\ell} L(y_i, z_i) \Big|_{z_i=a_{N-1}(x_i)}$$

$$\sum_{i=1}^{\ell} L(y_i, a_{N-1}(x_i) + b(x_i)) \rightarrow \min_b$$

- the direction calculated by taking into account the second derivatives of the loss function.
- penalties are added for the number of leaves
- criterion of informativeness, dependent on the optimal displacement vector.
- the stop criteria in the training of the tree depends on the shift

# Technical side: training in parallel

Which of the ensembling methods could be parallelized?

# Technical side: training in parallel

Which of the ensembling methods could be parallelized?

- Random Forest: parallel on the forest level (all trees are independent)

# Technical side: training in parallel

Which of the ensembling methods could be parallelized?

- Random Forest: parallel on the forest level (all trees are independent)
- Gradient boosting: parallel on one tree level

# Recap: ensembling methods

1. Bagging.
2. Random subspace method (RSM).
3. Bagging + RSM + Decision trees = Random Forest.
4. Gradient boosting.
5. Stacking.
6. Blending.

Great demo: [http://arogozhnikov.github.io/2016/06/24/gradient\\_boosting\\_explained.html](http://arogozhnikov.github.io/2016/06/24/gradient_boosting_explained.html)

# Recap: ensembling methods

1. Bagging.
2. Random subspace method (RSM).
3. Bagging + RSM + Decision trees = Random Forest.
4. Gradient boosting.
5. Stacking.
6. Blending.

Great thanks for materials to Radoslav Neychev

Great demo: [http://arogozhnikov.github.io/2016/06/24/gradient\\_boosting\\_explained.html](http://arogozhnikov.github.io/2016/06/24/gradient_boosting_explained.html)

# Links

1. Detailed description of bootstrapping procedure: [link](#)
2. Dyakonov blogpost about Gini coefficient and Gini Impurity: [link](#)
3. ODS ML course lesson about kNN and Decision Trees: [link](#)
4. Habr post about entropy in trees: [link](#)
5. Simply about bootstrapping on Habr: [link](#)
6. Notes about Decision Trees by Evgeny Sokolov: [link](#)